

Deep Learning : MP1 Report

Antoine Lavault

February 8, 2018

Getting Started

In this report, we will present the methods implemented for the first lab course of the Deep Learning class.

A Jupyter notebook for this mini-project can be found at : <https://github.com/ALavault/DL>

1 Simple Classification

The model here is the simplest possible : a single dense layer with a softmax activation

Layer (type)	Output Shape	Param #
dense_2 (Dense)	(None, 3)	15555
Total params: 15,555		
Trainable params: 15,555		
Non-trainable params: 0		

Figure 1: Description of single layer classifier

And the following code for obtaining such classifier :

```
1 from keras.models import Sequential
2 from keras import utils
3 model = Sequential()
4
5 from keras.layers import Dense, Activation, Flatten
6
7 model.add(Dense(3, input_shape=(X_train.shape[1],), activation='softmax'))
8
```

```

9  from keras.optimizers import SGD, adam
10 sgd = SGD(lr=0.01,decay=1e-6, momentum=0.9, nesterov=True)
11 adam = adam(lr=0.001, beta_1=0.9, beta_2=0.999, epsilon=None, decay=0.0, amsgrad=False)
12 model.compile(loss='categorical_crossentropy', optimizer=adam)
13
14
15 model.fit(X_train, Y_train, epochs=100, batch_size=32)

```

Here we choose the Adam method instead of the Stochastic Gradient Descent since the loss decreases faster with the Adam optimizer.

Moreover, the activation function for the layer is the softmax function. This activation function squashes the output in such a way that it will approximate the probability of being into a certain class. Since in its expression are intervening exponentials, it is also differentiable. As such, it can be trained.

Finally, we use 3 neurons because there are 3 different classes (square, circle, triangle) to classify. It's compulsory to have the same number of class as the training vector on the last layer (here the only one).

Just below, we present an example of the model prediction.

```

1  # Input
2  X_test = generate_a_disk()
3  X_test = X_test.reshape(1, X_test.shape[0])/255
4  model.predict(X_test)
5  # Output
6  array([[0., 1., 0.]], dtype=float32)

```

2 Visualization of the solution

To obtain these feature maps, we used the following script

```

1  weights = model.get_weights()
2  weightMat = np.asarray(weights[0]) # Obtaining the weights
3
4  # Plotting
5
6  plt.figure();
7  plt.subplot(1, 3, 1)
8  plt.imshow(weightMat[:,0].reshape(72,72), cmap='gray')
9  plt.title('Square')
10 plt.subplot(1, 3, 2)
11 plt.imshow(weightMat[:,1].reshape(72,72), cmap='gray')
12 plt.title('Circle')
13

```

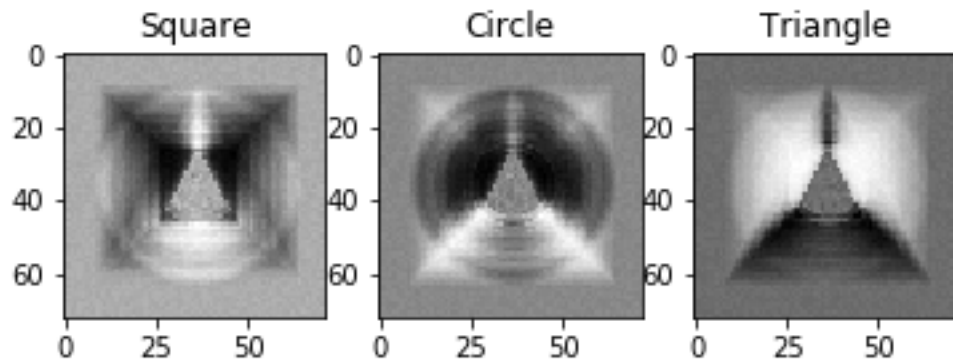


Figure 2: Features visualization

```

14 plt.subplot(1, 3, 3)
15 plt.imshow(weightMat[:,2].reshape(72,72), cmap='gray')
16 plt.title('Triangle')
17
18 plt.savefig('figures/part2_features')
19
20 plt.show()

```

The feature maps indeed show that the different shapes are learnt : the square map looks like a square, same for the circle and triangle.

In other words, we are visualizing how the network learnt how to distinguish between these different shapes all by itself (with a little help from a computer though).

3 A More Difficult Classification Problem

We will present 3 different networks : a linear classifier, a shallow CNN and a less-shallow CNN.

3.1 Linear classifier

We are using the same network as in the previous part.

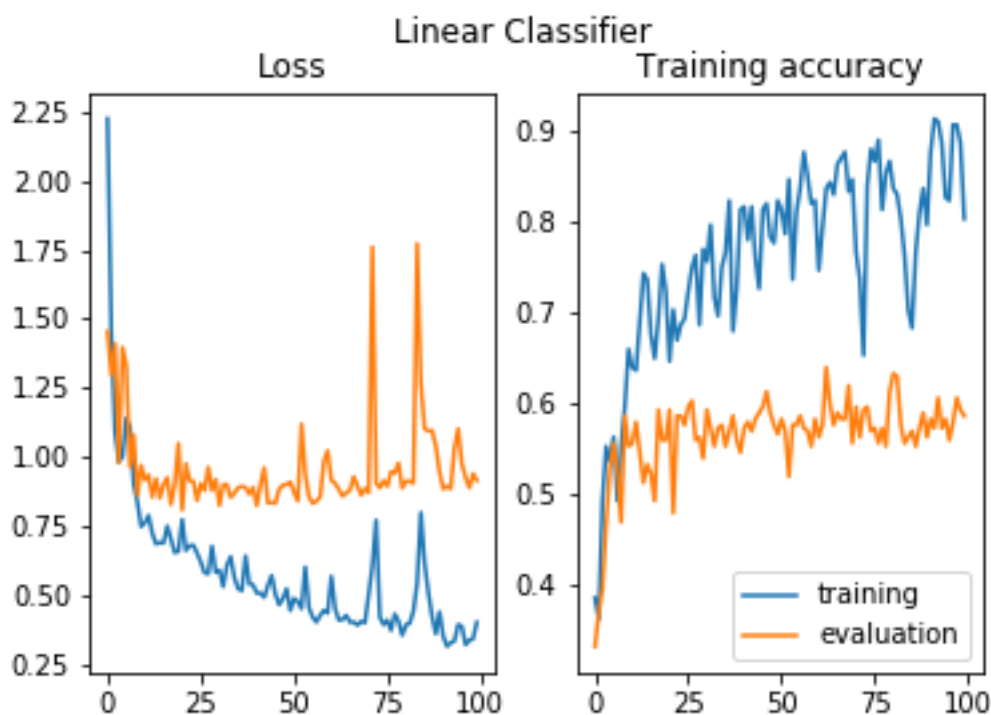


Figure 3: Loss and accuracy for a linear classifier

The performances are not exceptionnal : up to 70% accuracy on the training dataset and about 50% on the test dataset.

So, when it comes to "dynamic" images, the linear classifier is mostly lost.

3.2 Suggested CNN

The suggested CNN is implemented below :

```
1 from keras.models import Sequential
2 from keras import utils
```

```

3 model3 = Sequential()
4
5 from keras.layers import Dense, Activation, Flatten, Conv2D, MaxPooling2D
6 from keras.callbacks import History
7 history = History()
8
9 model3.add(Conv2D(16, (5,5), input_shape=(72,72,1) ) )
10 model3.add( MaxPooling2D() )
11 model3.add(Flatten())
12 model3.add(Dense(3, input_shape=(X_train.shape[1],), activation='softmax'))
13
14
15 from keras.optimizers import SGD, adam
16 sgd = SGD(lr=0.01,decay=1e-6, momentum=0.9, nesterov=True)
17 adam = adam(lr=0.001, beta_1=0.9, beta_2=0.999, epsilon=None, decay=0.0, amsgrad=False)
18
19 model3.compile(loss='categorical_crossentropy', optimizer=adam, metrics=['accuracy'])
20
21
22 model3.fit(X_train.reshape(X_train.shape[0],72,72,1), Y_train,
23 epochs=100, batch_size=32, callbacks=[history],
24 validation_data = (X_test.reshape(X_test.shape[0],72,72,1), Y_test),)

```

We can give a few comments. First, the performance is better than the linear classifier : we can get up to 70% on the test data. But we can also see from figure 4 that the model is overfitting after epoch 25.

Visually, we can say so because the loss function on the training data is decreasing when the loss on the testing data is increasing.

3.3 Another CNN

We then suggest another structure, a bit deeper than the last one.

```

1 from keras.models import Sequential
2 from keras import utils
3 model31 = Sequential()
4
5 from keras.layers import *
6 from keras.callbacks import History
7 history6 = History()
8
9 model31.add(Conv2D(16, (5,5), input_shape=(72,72,1) ) )
10 model31.add( MaxPooling2D() )
11 model31.add(Conv2D(16, (5,5) ) )
12 model31.add(BatchNormalization())

```




Figure 4: Loss and accuracy for the suggested CNN

```

13 model31.add(GlobalMaxPooling2D())
14
15
16 model31.add(Dense(16, activation='linear'))
17 model31.add( Dropout(0.4)) # To avoid overfitting
18 model31.add(Dense(3,activation='softmax'))
19
20 model31.compile(loss='categorical_crossentropy', optimizer='adam', metrics=['accuracy'])
21
22
23 model31.fit(X_train.reshape(X_train.shape[0],72,72,1), Y_train,
24 epochs=100, batch_size=32,
25 callbacks=[history6],
26 validation_data = (X_test.reshape(X_test.shape[0],72,72,1), Y_test),)

```

First thing we can see : it's the most accurate from all three. In fact, the accuracy can go as high as 96%.

By going deeper, we get more features and more features may give more




Figure 5: Loss and accuracy for a deeper CNN

informations. Also adding a dense layer between the output the convolutional/pooling layers gives a better accuracy to the model.

The batch normalization has been added since it gave better results when added in the model as well as the GlobalMaxPooling2D layer (which in fact gave the best performance increase).

Remark To obtain the different plots, we use the following script :

```

1 plt.figure(); plt.suptitle('Title...');
2
3 plt.subplot(1, 2, 1)
4 plt.plot(history.history['loss'])
5 plt.plot(history.history['val_loss'])
6 plt.title('Loss')
7
8 plt.subplot(1, 2, 2)
9 plt.plot(history.history['acc'])

```

```

10 plt.plot(history.history['val_acc'])
11 plt.legend(['training', 'evaluation'])
12 plt.title('Training accuracy')
13
14 plt.show()

```

4 A Regression Problem

From classification to regression... The main difference here is that we are going from a classification task, inducing discrete values to a regression task inducing continuous values.

Preprocessing A preprocessing which increased performances was just sorting the values (by axis). The following code does just that :

```

1 Y_train2 = Y_train.reshape(Y_train.shape[0], 3,2) # From vector to list of 'points'
2 y = Y_train2[0]
3 for i in range(len(Y_train2)):
4     y = Y_train2[i]
5     Y_train2[i] = y[y[:,0].argsort()]
6 Y_train2 = Y_train2.reshape(Y_train.shape[0],6) # Back to vector

```

Network The best network (whose performances are not extraordinary) is the following :

```

1 from keras.models import Sequential
2 from keras import utils
3 model4 = Sequential()
4
5 from keras.layers import *
6
7 model4.add(Conv2D(64, (7,7), input_shape=(72,72,1) ) )
8 model4.add( MaxPooling2D() )
9
10 model4.add(Conv2D(256, (5,5) ) )
11 model4.add(Conv2D(256, (5,5) ) )
12
13
14 model4.add( MaxPooling2D() )
15 model4.add(Conv2D(256, (5,5) ) )
16
17 model4.add( MaxPooling2D() )
18
19 model4.add(Conv2D(1024, (3,3) ) )

```



```

20 model4.add( MaxPooling2D() )
21
22 model4.add(BatchNormalization())
23 #model4.add(GlobalMaxPooling2D())
24
25
26
27 model4.add(Flatten())
28 model4.add(Dense(1024 , activation = 'sigmoid'))
29 model4.add( Dropout(0.5) )
30
31 model4.add(Dense(1024 , activation = 'sigmoid'))
32 model4.add( Dropout(0.5) )
33
34
35 model4.add(Dense(6, activation = 'sigmoid'))
36
37
38 from keras.optimizers import SGD, adam
39 sgd = SGD(lr=0.01,decay=1e-6, momentum=0.9, nesterov=True)
40 adam = adam(lr=0.001, beta_1=0.9, beta_2=0.999, epsilon=None, decay=0.0, amsgrad=False)
41 model4.compile(loss='mean_squared_error', optimizer='adam')
42
43
44 model4.fit(X_train, Y_train2, shuffle = True, epochs=1000, batch_size=16)

```

Such bloated network can at most approximate correctly on the training data but in this case, it is due to overfitting.

Also, since it is a regression problem, we need to change the loss function : the cross entropy is not meaningful in this case so we use the MSE (because it is simple).

5 An hourglass network

Basically, we will describe in this part an autoencoder.

First, the code :

```

1 autoencoder = Sequential()
2 from keras.layers import Dense, Activation, Flatten, Conv2D, MaxPooling2D,
3 UpSampling2D, BatchNormalization
4 from keras.callbacks import History
5 history5 = History()
6 neurons = 32
7
8 # Encode

```

```

9  autoencoder.add( Conv2D(neurons, (3, 3), input_shape = (72,72,1),
10 activation='relu', padding='same'))
11 autoencoder.add( MaxPooling2D((2, 2)))
12 autoencoder.add( Conv2D(neurons, (3, 3), activation='relu', padding='same') )
13
14 autoencoder.add(MaxPooling2D((2, 2)) )
15 autoencoder.add( Conv2D(neurons, (3, 3), activation='relu', padding='same') )
16
17 autoencoder.add(MaxPooling2D((2, 2)) )
18 autoencoder.add( Conv2D(neurons, (3, 3), activation='relu', padding='same') )
19 autoencoder.add( BatchNormalization() )
20
21 # Decode
22 autoencoder.add( UpSampling2D((2, 2)) )
23 autoencoder.add( Conv2D(neurons, (3, 3), activation='relu', padding='same') )
24 autoencoder.add( UpSampling2D((2, 2)) )
25 autoencoder.add( Conv2D(neurons, (3, 3), activation='relu', padding='same') )
26 autoencoder.add( UpSampling2D((2, 2)) )
27 autoencoder.add(Conv2D(1, (3, 3), activation='sigmoid', padding='same') )
28
29 autoencoder.compile(optimizer='adam', loss='binary_crossentropy')
30 import tensorflow as tf
31 # your code here
32 with tf.device('/gpu:0'):
33 autoencoder.fit(XX_train, XX_true, epochs=100, batch_size=128,
34 shuffle = True, callbacks=[history5], validation_data = (XX_test, XX_test_true) )

```

The use of the option "same" in padding allows the output to have the same size as the input. Thus, the hourglass structure will be easier to build (to have input and output layers with the same size).

Here we use the binary cross entropy because we want to measure the difference in terms of content, not absolute value.




Figure 6: Noisy examples

On the figure Before/After, we can see that noise is vanishing and shapes are kept.

As such, we can say that it indeed works.



Figure 7: Denoized examples

A remark on computing performance Keras, by using Tensorflow or Theano as a backend, can use either CPUs or GPUs. Note that it is compulsory to have an nVidia GPU since both backends do not support OpenCL.

As far as performances are concerned, using a GPU can lead to great speedups : for instance, on the networks presented in this report, the speedup factor can go from 70 to more than 250 between an intel i5 CPU and a nVidia Tesla GPU.

So, it is faster to train on a GPU. Just a limit, the memory of the GPU : sometimes it will be necessary to reduce the batch size to allow the computation to start (RessourceExhaustion error with TensorFlow).